

Introduction

Source: `vignettes/Introduction.Rmd` (<https://github.com/r-lib/R6/blob/master/vignettes/Introduction.Rmd>)

The R6 package provides a type of class which is similar to R's standard reference classes, but it is more efficient and doesn't depend on S4 classes and the methods package.

R6 classes

R6 classes are similar to R's standard reference classes, but are lighter weight, and avoid some issues that come along with using S4 classes (R's reference classes are based on S4). For more information about speed and memory footprint, see the Performance vignette.

Unlike many objects in R, instances (objects) of R6 classes have reference semantics. R6 classes also support:

- public and private methods
- active bindings
- inheritance (superclasses) which works across packages

Why the name R6? When R's reference classes were introduced, some users, following the names of R's existing class systems S3 and S4, called the new class system R5 in jest. Although reference classes are not actually called R5, the name of this package and its classes takes inspiration from that name.

The name R5 was also a code-name used for a different object system started by Simon Urbanek, meant to solve some issues with S4 relating to syntax and performance. However, the R5 branch was shelved after a little development, and it was never released.

Basics

Here's how to create a simple R6 class. The `public` argument is a list of items, which can be functions and fields (non-functions). Functions will be used as methods.

```
library(R6)

Person <- R6Class (../reference/R6Class.html)("Person",
  public = list(
    name = NULL,
    hair = NULL,
    initialize = function(name = NA, hair = NA) {
      self$name <- name
      self$hair <- hair
      self$greet()
    },
    set_hair = function(val) {
      self$hair <- val
    },
    greet = function() {
      cat(paste0("Hello, my name is ", self$name, ".\n"))
    }
  )
)
```

To instantiate an object of this class, use `$new()` :

```
ann <- Person$new("Ann", "black")
#> Hello, my name is Ann.
ann
#> <Person>
#>   Public:
#>     clone: function (deep = FALSE)
#>     greet: function ()
#>     hair: black
#>     initialize: function (name = NA, hair = NA)
#>     name: Ann
#>     set_hair: function (val)
```

The `$new()` method creates the object and calls the `initialize()` method, if it exists.

Inside methods of the class, `self` refers to the object. Public members of the object (all you've seen so far) are accessed with `self$x`, and assignment is done with `self$x <- y`.

Once the object is instantiated, you can access values and methods with `$` :

```
ann$hair
#> [1] "black"
ann$greet()
#> Hello, my name is Ann.
ann$set_hair("red")
ann$hair
#> [1] "red"
```

Implementation note: The external face of an R6 object is basically an environment with the public members in it. This is also known as the `public` environment. An R6 object's methods have a separate `private` environment which, roughly speaking, is the environment they “run in”. This is where `self` binding is found, and it is simply a reference back to public environment.

Private members

In the previous example, all the members were public. It's also possible to add private members:

```
Queue <- R6Class (../reference/R6Class.html)("Queue",
  public = list(
    initialize = function(...) {
      for (item in list(...)) {
        self$add(item)
      }
    },
    add = function(x) {
      private$queue <- c(private$queue, list(x))
      invisible(self)
    },
    remove = function() {
      if (private$length() == 0) return(NULL)
      # Can use private$queue for explicit access
      head <- private$queue[[1]]
      private$queue <- private$queue[-1]
      head
    }
  ),
  private = list(
    queue = list(),
    length = function() base::length (http://www.rdocumentation.org/packages/base/topic
  )
)

q <- Queue$new(5, 6, "foo")
```

Whereas public members are accessed with `self`, like `self$add()`, private members are accessed with `private`, like `private$queue`.

The public members can be accessed as usual:

```
# Add and remove items
q$add("something")
q$add("another thing")
q$add(17)
q$remove()
#> [1] 5
q$remove()
#> [1] 6
```

However, private members can't be accessed directly:

```
q$queue
#> NULL
q$length()
#> Error: attempt to apply non-function
```

A useful design pattern is for methods to return `self` (invisibly) when possible, because it makes them chainable. For example, the `add()` method returns `self` so you can chain them together:

```
q$add(10)$add(11)$add(12)
```

On the other hand, `remove()` returns the value removed, so it's not chainable:

```
q$remove()
#> [1] "foo"
q$remove()
#> [1] "something"
q$remove()
#> [1] "another thing"
q$remove()
#> [1] 17
```

Active bindings

Active bindings look like fields, but each time they are accessed, they call a function. They are always publicly visible.

```
Numbers <- R6Class (../reference/R6Class.html)("Numbers",
  public = list(
    x = 100
  ),
  active = list(
    x2 = function(value) {
      if (missing(value)) return(self$x * 2)
      else self$x <- value/2
    },
    rand = function() rnorm(1)
  )
)

n <- Numbers$new()
n$x
#> [1] 100
```

When an active binding is accessed as if reading a value, it calls the function with `value` as a missing argument:

```
n$x2
#> [1] 200
```

When it's accessed as if assigning a value, it uses the assignment value as the `value` argument:

```
n$x2 <- 1000
n$x
#> [1] 500
```

If the function takes no arguments, it's not possible to use it with `<-` :

```
n$rand
#> [1] 0.2648
n$rand
#> [1] 2.171
n$rand <- 3
#> Error: unused argument (quote(3))
```

Implementation note: Active bindings are bound in the public environment. The enclosing environment for these functions is also the public environment.

Inheritance

One R6 class can inherit from another. In other words, you can have super- and sub-classes.

Subclasses can have additional methods, and they can also have methods that override the superclass methods. In this example of a queue that retains its history, we'll add a `show()` method and override the `remove()` method:

```
# Note that this isn't very efficient - it's just for illustrating inheritance.
HistoryQueue <- R6Class (../reference/R6Class.html)("HistoryQueue",
  inherit = Queue,
  public = list(
    show = function() {
      cat("Next item is at index", private$head_idx + 1, "\n")
      for (i in seq_along(private$queue)) {
        cat(i, ": ", private$queue[[i]], "\n", sep = "")
      }
    },
    remove = function() {
      if (private$length() - private$head_idx == 0) return(NULL)
      private$head_idx <-< private$head_idx + 1
      private$queue[[private$head_idx]]
    }
  ),
  private = list(
    head_idx = 0
  )
)

hq <- HistoryQueue$new(5, 6, "foo")
hq$show()
#> Next item is at index 1
#> 1: 5
#> 2: 6
#> 3: foo
hq$remove()
#> [1] 5
hq$show()
#> Next item is at index 2
#> 1: 5
#> 2: 6
#> 3: foo
hq$remove()
#> [1] 6
```

Superclass methods can be called with `super$xx()`. The `CountingQueue` (example below) keeps a count of the total number of objects that have ever been added to the queue. It does this by overriding the `add()` method – it increments a counter and then calls the superclass's `add()` method, with `super$add(x)`:

```

CountingQueue <- R6Class (../reference/R6Class.html)("CountingQueue",
  inherit = Queue,
  public = list(
    add = function(x) {
      private$total <-< private$total + 1
      super$add(x)
    },
    get_total = function() private$total
  ),
  private = list(
    total = 0
  )
)

cq <- CountingQueue$new("x", "y")
cq$get_total()
#> [1] 2
cq$add("z")
cq$remove()
#> [1] "x"
cq$remove()
#> [1] "y"
cq$get_total()
#> [1] 3

```

Fields containing reference objects

If your R6 class contains any fields that also have reference semantics (e.g., other R6 objects, and environments), those fields should be populated in the `initialize` method. If the field is set to the reference object directly in the class definition, that object will be shared across all instances of the R6 objects. Here's an example:

```

SimpleClass <- R6Class (../reference/R6Class.html)("SimpleClass",
  public = list(x = NULL)
)

SharedField <- R6Class (../reference/R6Class.html)("SharedField",
  public = list(
    e = SimpleClass$new()
  )
)

s1 <- SharedField$new()
s1$e$x <- 1

s2 <- SharedField$new()
s2$e$x <- 2

# Changing s2$e$x has changed the value of s1$e$x
s1$e$x
#> [1] 2

```

To avoid this, populate the field in the `initialize` method:

```

NonSharedField <- R6Class (../reference/R6Class.html)("NonSharedField",
  public = list(
    e = NULL,
    initialize = function() self$e <- SimpleClass$new()
  )
)

n1 <- NonSharedField$new()
n1$e$x <- 1

n2 <- NonSharedField$new()
n2$e$x <- 2

# n2$e$x does not affect n1$e$x
n1$e$x
#> [1] 1

```

Other topics

Adding members to an existing class

It is sometimes useful to add members to a class after the class has already been created. This can be done using the `$set()` method on the generator object.

```

Simple <- R6Class (../reference/R6Class.html)("Simple",
  public = list(
    x = 1,
    getx = function() self$x
  )
)

Simple$set("public", "getx2", function() self$x*2)

# To replace an existing member, use overwrite=TRUE
Simple$set("public", "x", 10, overwrite = TRUE)

s <- Simple$new()
s$x
#> [1] 10
s$getx2()
#> [1] 20

```

The new members will be present only in instances that are created after `$set()` has been called.

To prevent modification of a class, you can use `lock_class=TRUE` when creating the class. You can also lock and unlock a class as follows:

```
# Create a locked class
Simple <- R6Class (../reference/R6Class.html)("Simple",
  public = list(
    x = 1,
    getx = function() self$x
  ),
  lock_class = TRUE
)

# This would result in an error
# Simple$set("public", "y", 2)

# Unlock the class
Simple$unlock()

# Now it works
Simple$set("public", "y", 2)

# Lock the class again
Simple$lock()
```

Cloning objects

By default, R6 objects have method named `clone` for making a copy of the object.

```
Simple <- R6Class (../reference/R6Class.html)("Simple",
  public = list(
    x = 1,
    getx = function() self$x
  )
)

s <- Simple$new()

# Create a clone
s1 <- s$clone()
# Modify it
s1$x <- 2
s1$getx()
#> [1] 2

# Original is unaffected by changes to the clone
s$getx()
#> [1] 1
```

If you don't want a `clone` method to be added, you can use `cloneable=FALSE` when creating the class. If any loaded R6 object has a `clone` method, that function uses 75.1 kB, but for each additional object, the `clone` method costs a trivial amount of space (112 bytes).

Deep cloning

If there are any fields which are objects with reference semantics (environments, R6 objects, reference class objects), the copy will get a reference to the same object. This is sometimes desirable, but often it is not.

For example, we'll create an object `c1` which contains another R6 object, `s`, and then clone it. Because the original's and the clone's `s` fields both refer to the same object, modifying it from one results in a change that is reflected in the other.

```
Simple <- R6Class (../reference/R6Class.html)("Simple", public = list(x = 1))

Cloneable <- R6Class (../reference/R6Class.html)("Cloneable",
  public = list(
    s = NULL,
    initialize = function() self$s <- Simple$new()
  )
)

c1 <- Cloneable$new()
c2 <- c1$clone()

# Change c1's `s` field
c1$s$x <- 2

# c2's `s` is the same object, so it reflects the change
c2$s$x
#> [1] 2
```

To make it so the clone receives a **gst** of `s`, we can use the `deep=TRUE` option:

```
c3 <- c1$clone(deep = TRUE)

# Change c1's `s` field
c1$s$x <- 3

# c2's `s` is different
c3$s$x
#> [1] 2
```

The default behavior of `clone(deep=TRUE)` is to copy fields which are R6 objects, but not copy fields which are environments, reference class objects, or other data structures which contain other reference-type objects (for example, a list with an R6 object).

If your R6 object contains these types of objects and you want to make a deep clone of them, you must provide your own function for deep cloning, in a private method named `deep_clone`. Below is an example of an R6 object with two fields, `a` and `b`, both of which are environments, and both of which contain a value `x`. It also has a field `v` which is a regular (non-reference) value, and a private `deep_clone` method.

The `deep_clone` method is called once for each field. It is passed the name and value of the field, and the value it returns is used in the clone.

```

CloneEnv <- R6Class (../reference/R6Class.html)("CloneEnv",
  public = list(
    a = NULL,
    b = NULL,
    v = 1,
    initialize = function() {
      self$a <- new.env(parent = emptyenv())
      self$b <- new.env(parent = emptyenv())
      self$a$x <- 1
      self$b$x <- 1
    }
  ),
  private = list(
    deep_clone = function(name, value) {
      # With x$clone(deep=TRUE) is called, the deep_clone gets invoked once for
      # each field, with the name and value.
      if (name == "a") {
        # `a` is an environment, so use this quick way of copying
        list2env(as.list.environment(value, all.names = TRUE),
          parent = emptyenv())
      } else {
        # For all other fields, just return the value
        value
      }
    }
  )
)

c1 <- CloneEnv$new()
c2 <- c1$clone(deep = TRUE)

```

When `c1$clone(deep=TRUE)` is called, the `deep_clone` method is called for each field in `c1`, and is passed the name of the field and value. In our version, the `a` environment gets copied, but `b` does not, nor does `v` (but that doesn't matter since `v` is not a reference object). We can test out the clone:

```

# Modifying c1$a doesn't affect c2$a, because they're separate objects
c1$a$x <- 2
c2$a$x
#> [1] 1

# Modifying c1$b does affect c2$b, because they're the same object
c1$b$x <- 3
c2$b$x
#> [1] 3

# Modifying c1$v doesn't affect c2$v, because they're not reference objects
c1$v <- 4
c2$v
#> [1] 1

```

In the example `deep_clone` method above, we checked the name of each field to determine what to do with it, but we could also check the value, by using `inherits(value, "R6")`, or `is.environment()`, and so on.

Printing R6 objects to the screen

R6 objects have a default `print` method that lists all members of the object. If a class defines a `print` method, then it overrides the default one.

```
PrettyCountingQueue <- R6Class (../reference/R6Class.html)("PrettyCountingQueue",
  inherit = CountingQueue,
  public = list(
    print = function(...) {
      cat("<PrettyCountingQueue> of ", self$get_total(), " elements\n", sep = "")
    }
  )
)
```

```
pq <- PrettyCountingQueue$new(1, 2, "foobar")
pq
#> <PrettyCountingQueue> of 3 elements
```

Finalizers

Sometimes it's useful to run a function when the object is garbage collected. For example, you may want to make sure a file or database connection gets closed. To do this, you can define a `finalize()` method, which will be called with no arguments when the object is garbage collected.

```
A <- R6Class (../reference/R6Class.html)("A", public = list(
  finalize = function() {
    print("Finalizer has been called!")
  }
))

# Instantiate an object:
obj <- A$new()

# Remove the single existing reference to it, and force garbage collection
# (normally garbage collection will happen automatically from time
# to time)
rm(obj); gc()
#> [1] "Finalizer has been called!"
#>          used (Mb) gc trigger (Mb) limit (Mb) max used (Mb)
#> Ncells  611730 32.7   1199437 64.1      NA   1199437 64.1
#> Vcells 1183901  9.1    8388608 64.0    16384   2402993 18.4
```

Finalizers are implemented using the `reg.finalizer()` function, and they set `onexit=TRUE`, so that the finalizer will also be called when R exits. This is useful in some cases, like database connections.

Class methods vs. member functions

When an R6 class definition contains functions in the public or private sections, those functions are **global** and they can access `self` (as well as `private` and `super` when available). When an R6 object is cloned, the resulting object's methods will have a `self` that refers to the new object. This works by changing the `environment` of the method in the cloned object.

In contrast to class methods, you can also add regular functions as members of an R6 object. This can be done by assigning the function to a field in the `initialize` method, or after the object has been instantiated. These functions are not class methods, and they will not have access to `self`, `private`, or `super`.

Here is a trivial class which has a method `get_self()` that simply returns `self`, as well as an empty member, `fn`. In this example, we'll assign a function to `fn` which has the same body as `get_self`. However, since it's a regular function, `self` will refer to something other than the R6 object:

```
FunctionWrapper <- R6Class (../reference/R6Class.html)("FunctionWrapper",
  public = list(
    get_self = function() {
      self
    },
    fn = NULL
  )
)

a <- FunctionWrapper$new()

# Create a function that accesses a variable named `self`.
# Note that `self` in this function's scope refers to 100, not to the R6 object.
self <- 100
a$fn <- function() {
  self
}

a$get_self()
#> <FunctionWrapper>
#>   Public:
#>   clone: function (deep = FALSE)
#>   fn: function ()
#>   get_self: function ()

a$fn()
#> [1] 100
```

As of R6 2.3.0, if the object is cloned, member (non-method) functions will not have their enclosing environment changed, which is what one would normally expect. It will behave this way:

```
b <- a$clone()

b$get_self()
#> <FunctionWrapper>
#>   Public:
#>   clone: function (deep = FALSE)
#>   fn: function ()
#>   get_self: function ()

b$fn()
#> [1] 100
```

